

Credit Card Fraud Transaction Detection

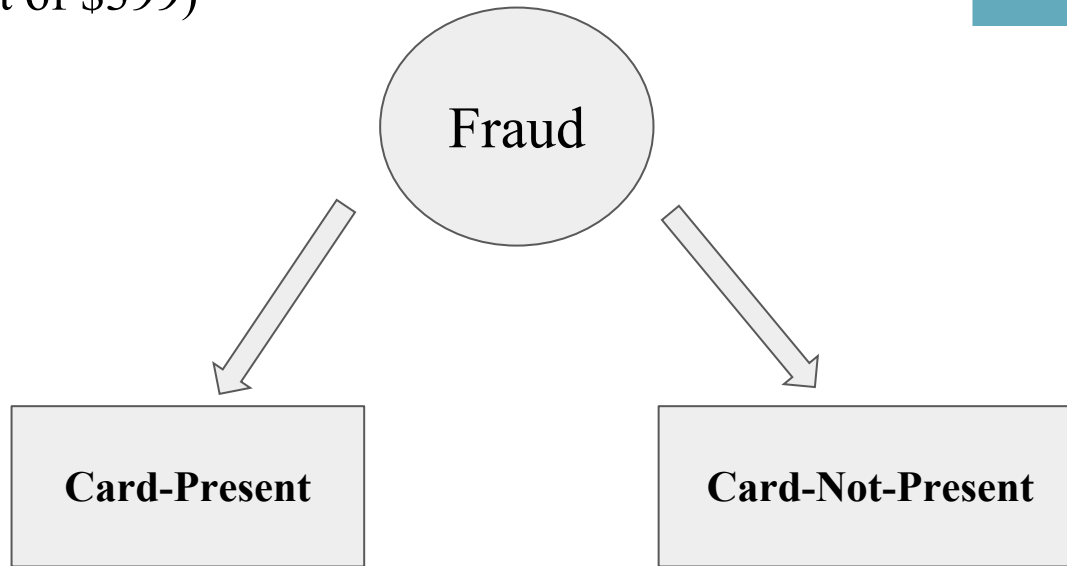
Module 2

Model Building – Optimization – Performance Evaluation – Tackling Imbalance Class

Group 6 Evan Day

Recap-Background

1/10 people in US has been a victim of credit card fraud
(median amount of \$399)



- Lost/Stolen
- Counterfeited card
- Card not received

- Data Leak
- Fraud Website

Recap-Background

Real World data can be **CONFIDENTIAL** : Privacy Issue

Transaction Data Frame:

TRANSACTION_ID, DATETIME, CUSTOMER_ID, Store_ID, AMOUNT, FRAUD

Customer Data Frame:

CUSTOMER_ID, location_x, location_y, Mean_trans, STD_trans, num_tran_day

Store Data Frame:

Store_ID, Station_location_x, Station_location_y, Mean_trans, STD_trans, num_tran_day

Feature Engineering of dataset (Improvement on Data set)

- Transaction during night, Transaction during weekend
- RFM Model (Recency, Frequency, Monetary Value) [Véronique, 2015]:
 - a. # of Transactions within a window [1, 7, 30]
 - b. Avg. amount within the window
- Risk Score: Avg. fraudulent transactions that occurred on a store [1, 7, 30]
- ...

Final Data Set

Python

...

TRANSACTION_ID

DATETIME

CUSTOMER_ID

STORE_ID

AMOUNT

Trans_TIME_SECONDS

Trans_TIME_DAYS

Trans_FRAUD

Trans_DURING_WEEKEND

Trans_DURING_NIGHT

...

CUS

0

0

2022-01-01 00:00:17

5820

2647

17.99

17

0

0

1

1

...

1

1

2022-01-01 00:00:17

6160

2980

44.48

17

0

0

1

1

...

2

2

2022-01-01 00:00:30

356

752

86.87

30

0

0

1

1

...

3

3

2022-01-01 00:00:31

1829

2266

16.65

31

0

0

1

1

...

4

5

2022-01-01 00:00:31

6820

8046

128.04

31

0

0

1

1

...

...

...

...

...

...

...

...

...

...

...

...

2421220

2421220

2022-04-10 23:58:31

3491

2534

144.45

8639911

99

0

1

0

...

2421221

2421221

2022-04-10 23:59:24

4381

7285

80.93

8639964

99

0

1

0

...

2421222

2421222

2022-04-10 23:59:27

3805

7758

39.65

8639967

99

0

1

0

...

2421223

2421223

2022-04-10 23:59:38

4491

8719

142.47

8639978

99

0

1

0

...

2421224

2421224

2022-04-10 23:59:48

9535

3388

38.94

8639988

99

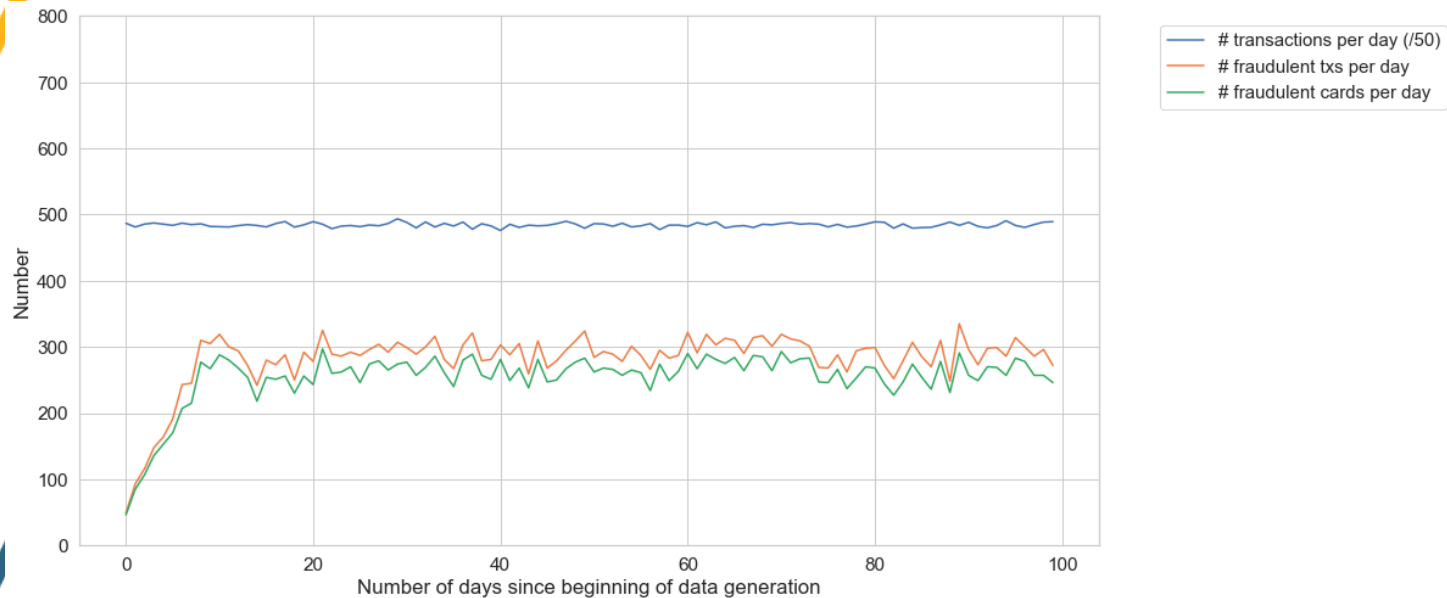
0

1

0


...

2421225 rows x 22 columns



How to Predict?

- It is important to note that we chose to run the test set one week after the last trade in the training set. In the context of fraud detection, this period of time that separates the training and test sets is called the *delay time or feedback delay* [Andrea, 2017]. It accounts for the fact that in real-world fraud detection systems, the label of a transaction, fraudulent or real, is known only after a customer complaint or as a result of a fraud investigation.



Model Building

- Decision Trees
- Logistic Regressions
- Random Forest
- XG-Boosting

Performance Metrics

- ROC/AUC:

The ROC curve is used to assess the overall diagnostic performance of a test and to compare the performance of two or more diagnostic tests. It is also used to select an optimal cut-off value for determining the presence or absence of a disease.

- Card Precision Top K (CP@K)

multiple fraudulent transactions from the same card should count as a single correct detection since investigators check all the recent transactions when contacting cardholders.



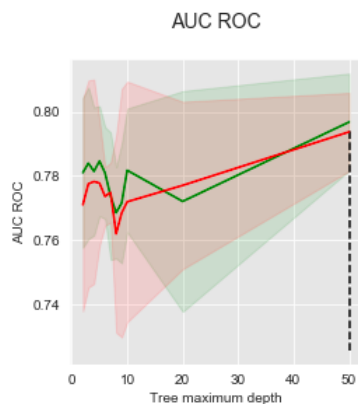
- Average Precision

(Why choose these? Instead of traditional Accuracy? Imbalance Class, explained later)

Model Building

- Decision Trees

	AUC ROC	Average precision	Card Precision@100
Best estimated param	50	5	5
Validation performance	0.794+/-0.01	0.548+/-0.02	0.38+/-0.02
Test performance	0.797+/-0.01	0.553+/-0.03	0.377+/-0.03
Optimal parameters	50	5	3
Optimal test performance	0.797+/-0.01	0.553+/-0.03	0.38+/-0.03




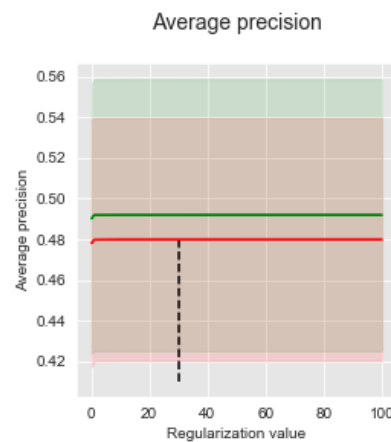
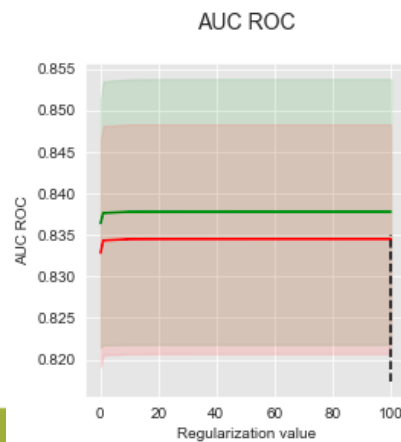
- Maximum depth are different according to different performance metrics
- The best parameters for the validation may not be the optimal parameter for the test set. CP@100.

Model Building

- Logistic Regressions → HP: regularization parameter

+

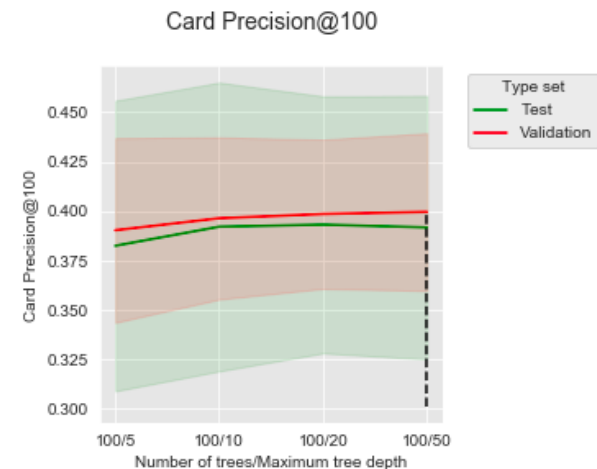
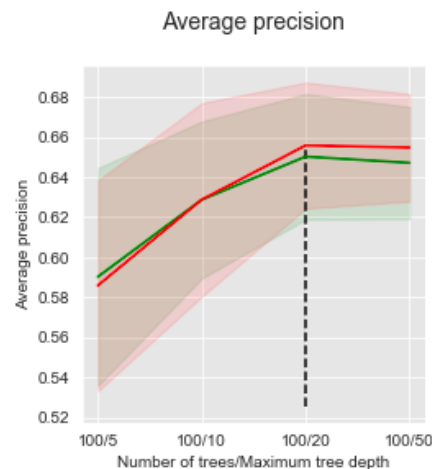
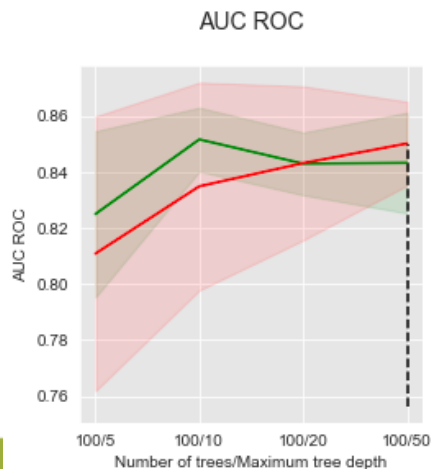
	AUC ROC	Average precision	Card Precision@100
Best estimated param	100.0	30.0	0.1
Validation performance	0.835+/-0.01	0.48+/-0.03	0.353+/-0.03
Test performance	0.838+/-0.01	0.492+/-0.03	0.361+/-0.03
Optimal parameters	100.0	100.0	0.1
Optimal test performance	0.838+/-0.01	0.492+/-0.03	0.361+/-0.03



Model Building

- Random Forest → HP: MAX Depth and # of trees



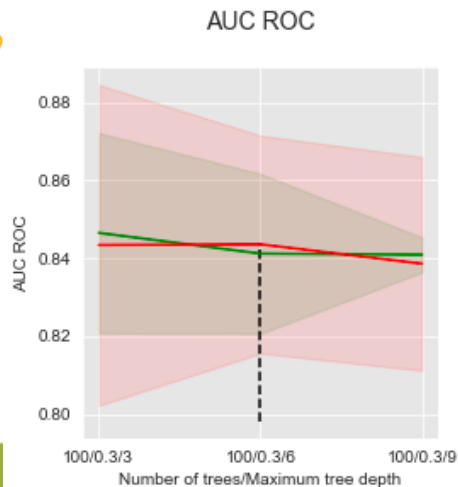
	AUC ROC	Average precision	Card Precision@100
Best estimated param	100/50	100/20	100/50
Validation performance	0.85+/-0.01	0.656+/-0.02	0.4+/-0.02
Test performance	0.843+/-0.01	0.65+/-0.02	0.392+/-0.03
Optimal parameters	100/10	100/20	100/20
Optimal test performance	0.852+/-0.01	0.65+/-0.02	0.393+/-0.03




Model Building

- XG Boosting → HP: MAX Depth, # of trees, learning rate

	AUC ROC	Average precision	Card Precision@100
Best estimated param	100/0.1/6	100/0.1/3	100/0.1/3
Validation performance	0.845+/-0.02	0.652+/-0.01	0.397+/-0.02
Test performance	0.848+/-0.01	0.644+/-0.01	0.389+/-0.03
Optimal parameters	100/0.1/3	100/0.3/3	100/0.3/3
Optimal test performance	0.853+/-0.01	0.647+/-0.02	0.39+/-0.04

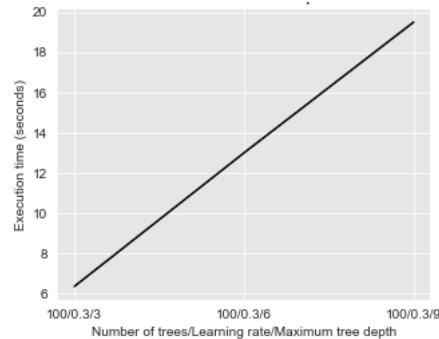



Model Selection:

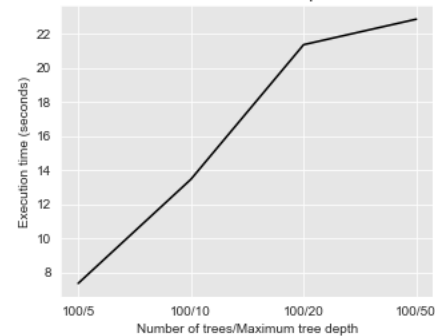


On this simulated dataset, XGBoost performs the best in terms of average accuracy and CP@100, followed by random forest, logistic regression, and finally decision tree. The gap in performance is most pronounced under the average precision metric. In terms of AUC ROC, the performance of logistic regression, random forest and XGBoost is very similar.

Model Selection - Time:



XG-Boosting



RF

It illustrates that model selection for complex models such as random forests or boosting usually requires more computational resources, since they require tuning a higher number of hyperparameters.

Further Optimization – Random Research:

```
print("Total execution time for XGBoost with grid search: "+str(round(execution_time_boosting,2))+ "s")  
print("Total execution time for XGBoost with random search: "+str(round(execution_time_boosting_random,2))+ "s")
```

[48] ✓ 0.0s

Python

```
... Total execution time for XGBoost with grid search: 1533.7s  
Total execution time for XGBoost with random search: 1212.4s
```


Imbalance Class Issue:

Only 3% fraud, the model can predict all the transaction are all the normal transaction, but we can still achieve an accuracy =97%! As a result, we should consider how to resolve the imbalance problem:

Using the **Cost Sensitive Learning: Used if the misclassification cost is not equal.**


For visualization purpose, we retrieve the Top-2 PCA to represent the X-variables.

We can see from the plot that the distribution of two classes are overlapping, however, we can see some patterns here.




Imbalance Class Issue:

Normal Decision Tree:




Adding the Cost-Sensitive learning




Imbalance Class Issue:

Normal Logistic Regression:




Adding the Cost-Sensitive learning




Improvement from Logistic Regression

Normal Logistic Regression:



Adding the Cost-Sensitive learning



Limitations

- Cost-sensitive learning is ambiguous, we cannot accurately define the “Cost”.
- Machine learning may not be the best choice because there are limited hyper-parameters we can control
- Data Limitation: Pure simulated data have many drawbacks
- ...

Further Action

- Resampling methods e.g. SMOTE
- Adding deep learning in our further work, building layers to feed forward
- Instead of focusing on the overall performance, we are trying to focus on fraud class accuracy, because it is cheap to say a normal transaction is fraud.
- ...



Q&A



Thanks

References

Charles X Ling and Victor S Sheng. Cost-sensitive learning and the class imbalance problem. *Encyclopedia of machine learning*, 2011:231–235, 2008.

Véronique Van Vlasselaer, Cristián Bravo, Olivier Caelen, Tina Eliassi-Rad, Leman Akoglu, Monique Snoeck, and Bart Baesens. Apate: a novel approach for automated credit card transaction fraud detection using network-based extensions. *Decision Support Systems*, 75:38–48, 2015.

Machine Learning Group: Reproducible Machine Learning for Credit Card Fraud detection - Practical handbook